

Transaction

- **Definition:** A sequence of database operations (like reads and writes) treated as a single, logical unit of work. It's "all or nothing."
- **Goal:** To ensure data consistency and integrity, especially when multiple operations are involved (e.g., transferring funds requires debiting one account and crediting another).
- **Key Operations:** COMMIT (make changes permanent), ROLLBACK (undo changes).

ACID Properties

- **Definition:** Fundamental properties guaranteeing reliable transaction processing. Crucial for maintaining data integrity.
- **Properties:**
 - **Atomicity:** Ensures transactions are indivisible. Either all operations within the transaction succeed, or none are applied. Prevents partial updates.
 - **Consistency:** Ensures a transaction transitions the database from one valid state to another, preserving all predefined rules and constraints (like foreign keys, unique constraints).
 - **Isolation:** Ensures that concurrent transactions do not interfere with each other inappropriately. The effect should be the same as if transactions ran sequentially. This prevents anomalies like dirty reads.
 - **Durability:** Ensures that once a transaction is committed, its changes are permanent and survive system failures (crashes, power outages). Often achieved using mechanisms like Write-Ahead Logging (WAL).

Concurrency Anomalies & Isolation Levels

- **Problem:** When multiple transactions run concurrently without sufficient isolation, anomalies can occur.
- **Common Anomalies:**
 - **Dirty Read:** Reading data that has been modified by another transaction but *not yet committed*. (Risk: Reading data that might be rolled back).
 - **Non-Repeatable Read:** Reading the same data twice within a transaction and getting different values because another transaction modified and committed the data in between. (Risk: Inconsistent view of data within a single transaction).
 - **Phantom Read:** Re-executing a query within a transaction returns a different set of rows because another transaction inserted/deleted rows matching the criteria and committed. (Risk: Set of data appears to change unexpectedly).
- **Isolation Levels:** Control the trade-off between consistency and concurrency by defining which anomalies are prevented. Common levels (from least to most strict):
 1. **Read Uncommitted:** Allows all anomalies. Rarely used.
 2. **Read Committed:** Prevents Dirty Reads. Default for many databases (e.g., PostgreSQL, SQL Server). Still allows Non-Repeatable and Phantom Reads.
 3. **Repeatable Read:** Prevents Dirty Reads and Non-Repeatable Reads. Still allows Phantom Reads. Default for MySQL (InnoDB).
 4. **Serializable:** Prevents all anomalies. Provides the highest level of isolation, effectively executing transactions as if they were serial. Can significantly impact concurrency.

Concurrency Control Mechanisms

- **Purpose:** Techniques used by DBMS to manage concurrent access and prevent anomalies, ensuring some level of serializability.
- **Lock-Based Protocols:** Common approach where transactions acquire locks on data before accessing it.
 - **Shared Lock (S-lock / Read Lock):** Allows multiple transactions to read concurrently.
 - **Exclusive Lock (X-lock / Write Lock):** Allows only one transaction to write (or read). Prevents other locks.
 - **Two-Phase Locking (2PL):** A protocol ensuring serializability (though not necessarily preventing deadlocks). Transactions have a "growing" phase (acquiring locks) and a "shrinking" phase (releasing locks). *Strict 2PL* (holding write locks until commit/abort) is common as it prevents cascading rollbacks.
 - **Deadlock:** A situation where two or more transactions are blocked indefinitely, each waiting for a lock held by the other. Databases have mechanisms to detect and resolve deadlocks (often by aborting one transaction).
- **Timestamp-Based Protocols:** Assigns timestamps to transactions and uses them to determine execution order, potentially aborting transactions that violate the order. Less common than locking in many traditional RDBMSs.
- **Optimistic Concurrency Control (OCC):** Assumes conflicts are rare. Transactions proceed without locking, and checks are performed at commit time. If a conflict is detected, the transaction rolls back.
- **Multiversion Concurrency Control (MVCC):** Allows reads and writes to proceed concurrently by maintaining multiple versions of data. Reads access a consistent snapshot without blocking writes. Used by databases like PostgreSQL and Oracle.

Schedules & Serializability

- **Schedule:** The order in which operations from concurrent transactions are executed.
- **Serial Schedule:** Transactions run one after another. Consistent but slow.
- **Concurrent Schedule:** Operations are interleaved. Faster but risks anomalies.
- **Serializability:** The goal for concurrent schedules. Ensures the outcome is equivalent to *some* serial execution, maintaining consistency. Conflict-serializability (based on ordering conflicting operations) is the most common type.

Recoverability

- **Recoverable:** Ensures transactions only commit *after* reading data from already committed transactions. Prevents committing based on "dirty" reads.
- **Cascadeless:** Prevents cascading rollbacks. Ensures transactions only read data from committed transactions. (Strict 2PL achieves this).

Database Recovery Management

- **Purpose:** Restore database consistency after failures.
- **Key Techniques:**
 - **Logging (Write-Ahead Logging - WAL):** Changes are written to a log file *before* being applied to the database itself. Essential for Undo (rollback) and Redo (durability).
 - **Backup and Restore:** Creating periodic copies of the database to restore from in case of major failures (like media failure).